# Kubernetes Fundmentals

.NET CORE

Kubernetes is a portable, extensible, open-source platform for managing containerized workloads and services. It facilitates both declarative configuration and automation.
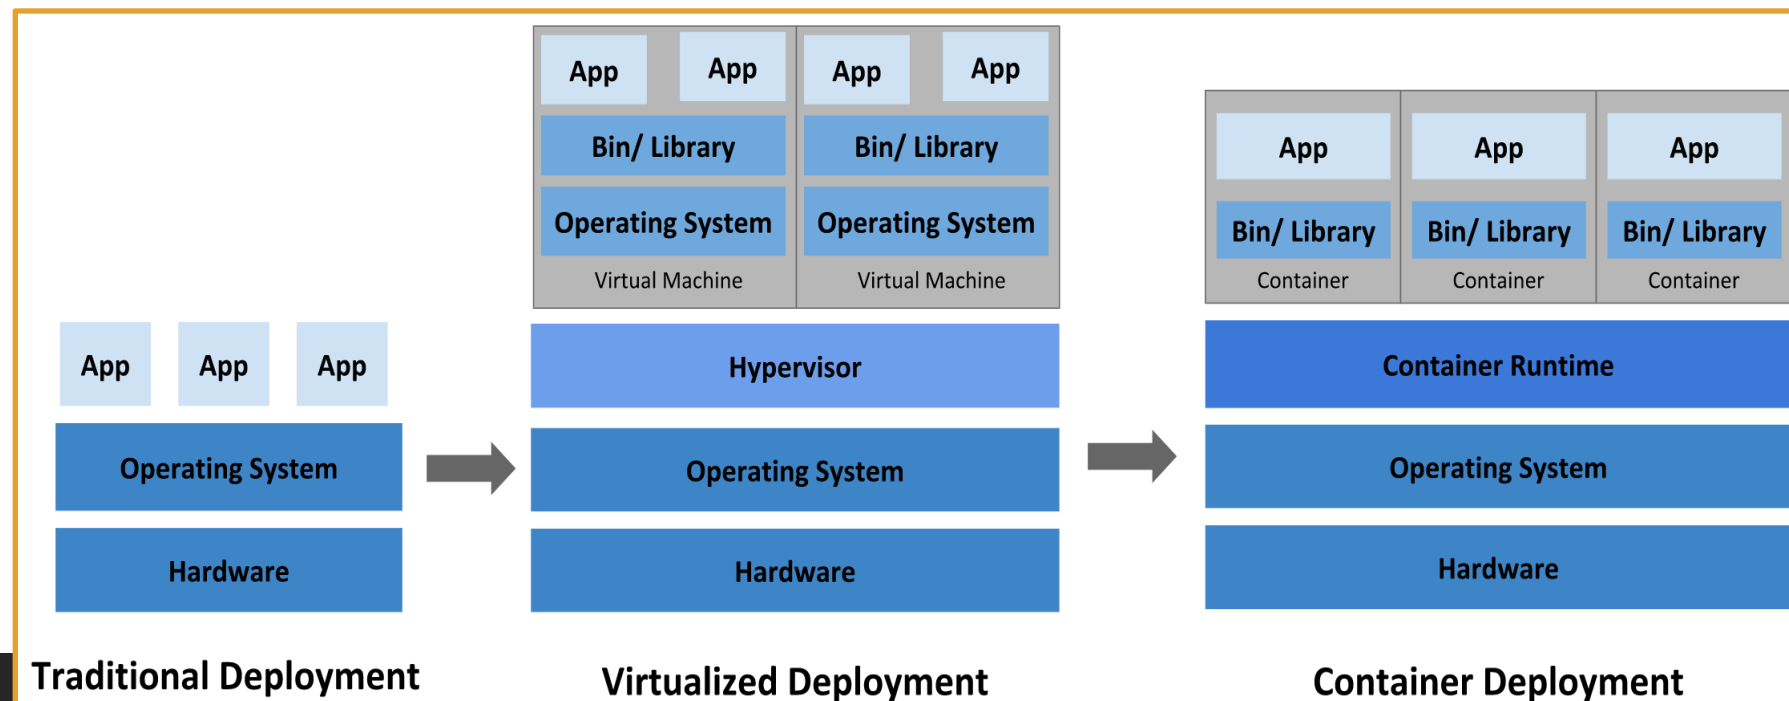
What is Kubernetes?

Kubernetes is a production-grade, open-source infrastructure for the deployment, scaling, management, and composition of application containers across *clusters* of hosts. It is inspired by previous work at Google *Kubernetes project*. The name *Kubernetes* originates from Greek, meaning helmsman or pilot.

Traditional Deployment **Virtualized Deployment** **Container Deployment**
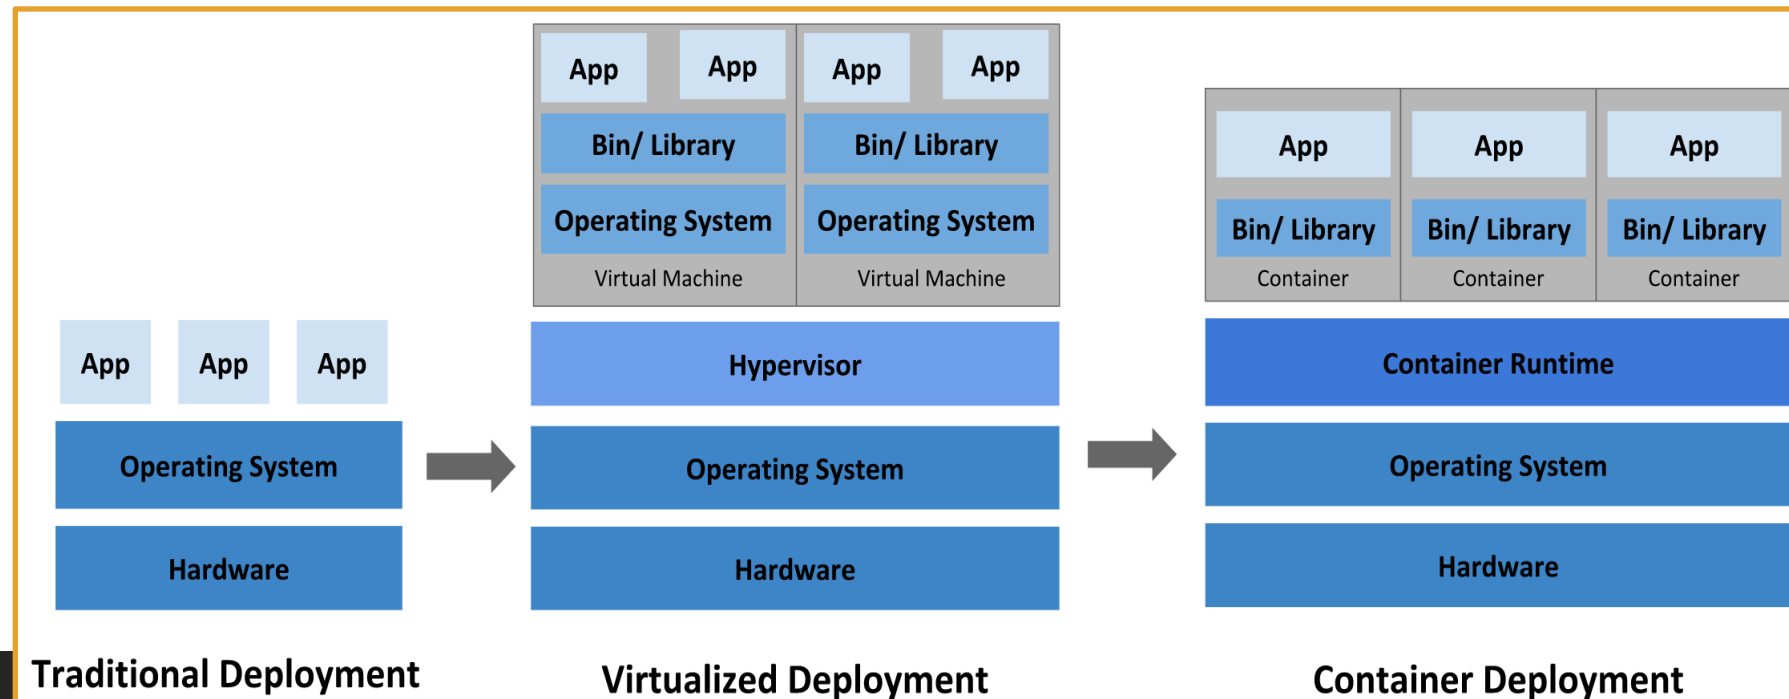
What is Kubernetes?

Kubernetes provides you with a framework to run distributed systems resiliently. It takes care of scaling and failover and provides deployment patterns. It allows you to automate the deployment of your containerized microservices. This makes it easier to manage all the components and microservices in your application.
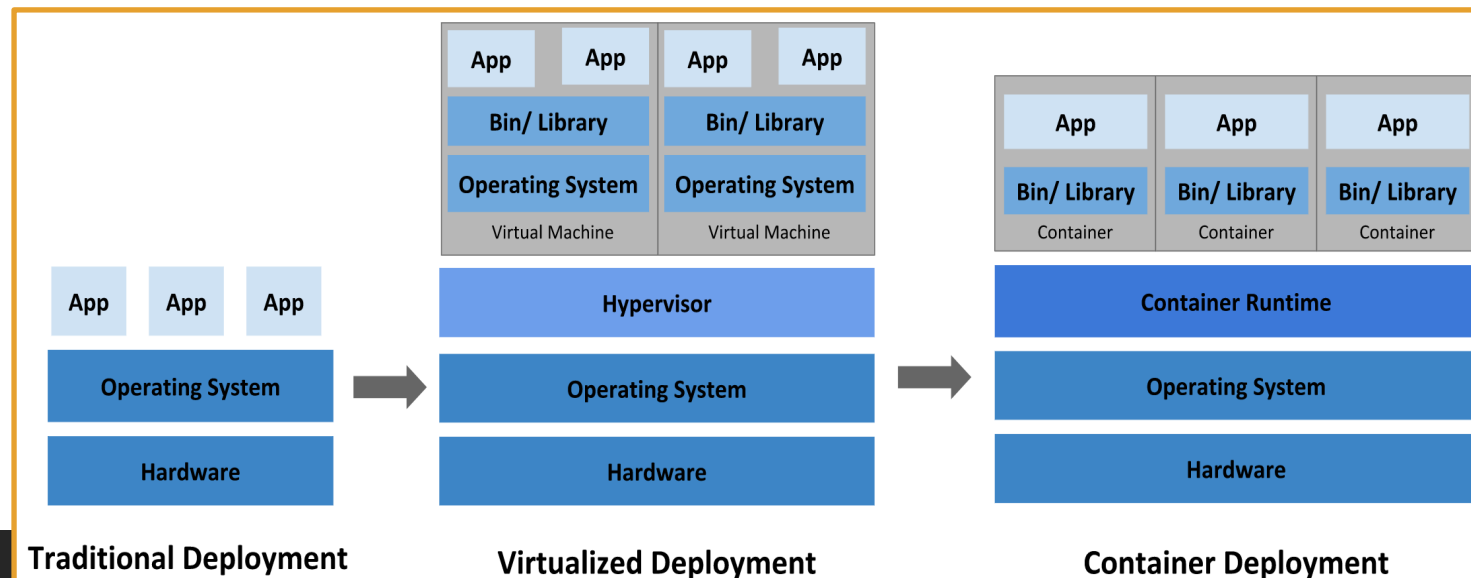

Traditional Deployment   **Virtualized Deployment**   **Container Deployment**

What is Kubernetes?

| Kubernetes containers allow you to:			
Deploy images quickly	Maintain CI/CD	Enhance Separation of Concerns	Run your application anywhere
Have an elastic, scalable MSA	Isolate resources	Use resources effectively	Run your application on any platform

Traditional Deployment **Virtualized Deployment** **Container Deployment**

Kubernetes Architecture – Overview (1/2)

https://kubernetes.io/docs/concepts/overview/components/

- *Kubernetes* is not a traditional, all-inclusive PaaS (Platform as a Service) system. *Kubernetes* operates at the container level rather than at the hardware level.
- When you deploy *Kubernetes*, you get a cluster.
- A Cluster consists of worker machines (nodes), that run *containerized* applications.

Kubernetes Architecture – Overview (2/2)

https://kubernetes.io/docs/concepts/overview/components/
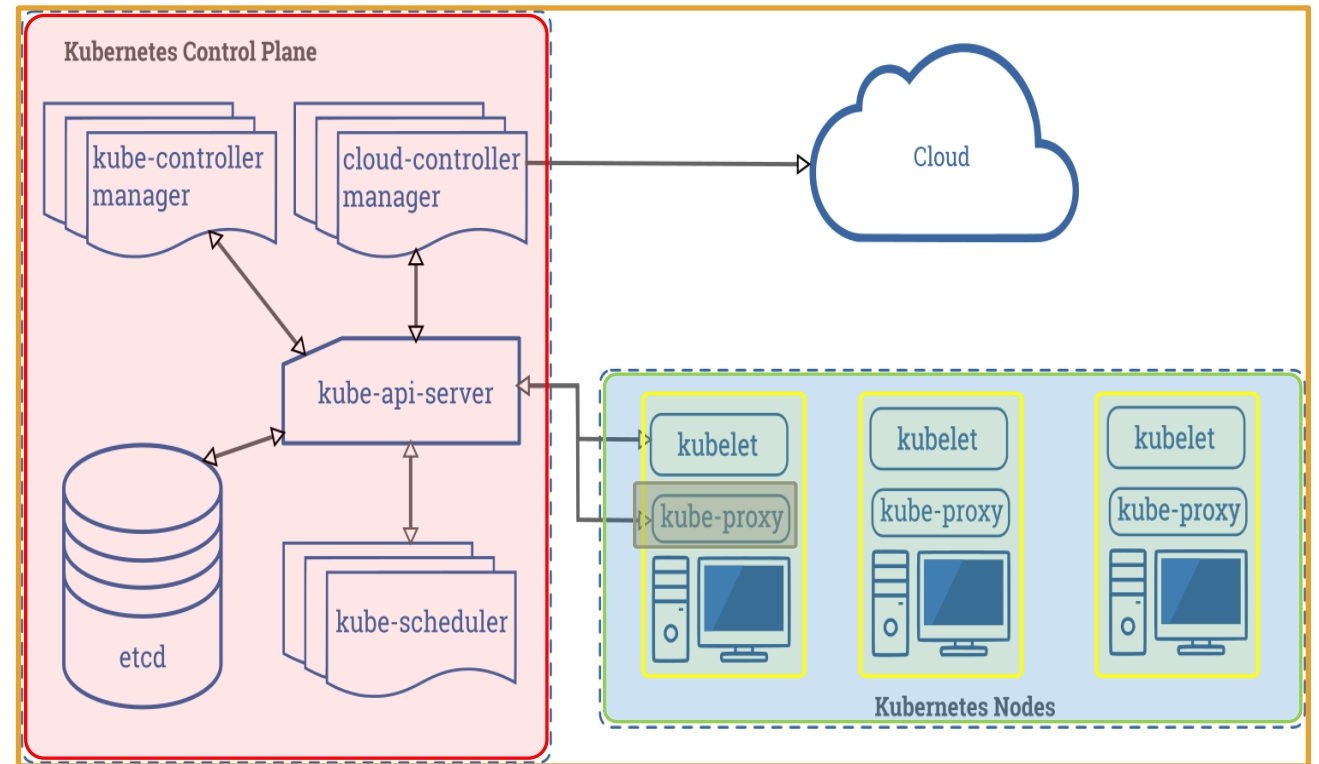
- The worker node(s) host the Pods that are the components of the application workload.
- The control plane manages the worker nodes and the Pods in the cluster.
- In production environments, the control plane usually operates across multiple computers and a cluster usually runs multiple nodes. This provides fault-tolerance and high availability.
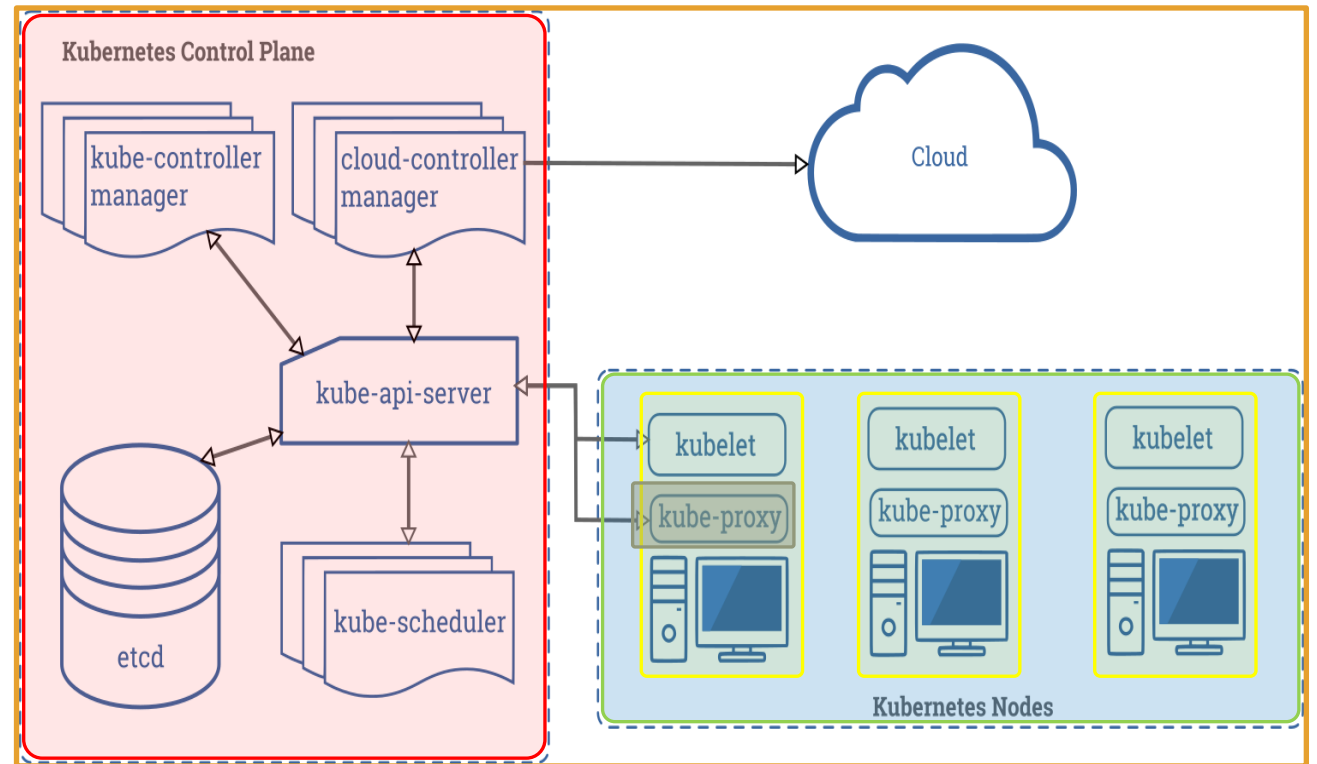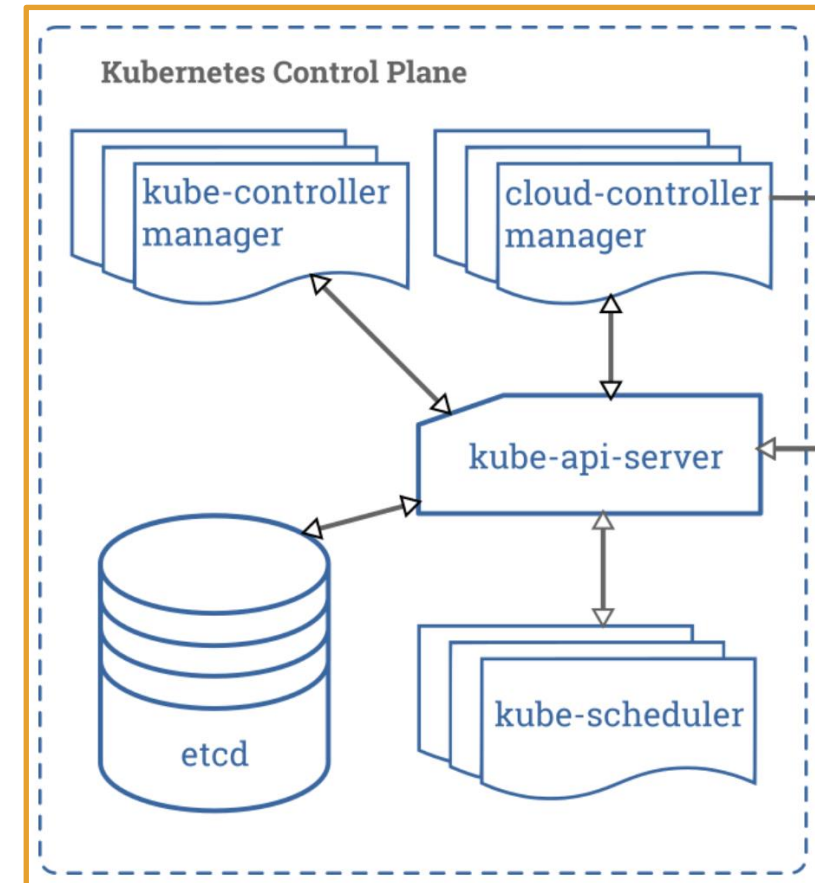
Kubernetes Control Plane (Master)

The *control plane's* components make global decisions about the cluster, as well as detecting and responding to *cluster* events (for example, starting up a new pod when a deployment's 'replicas' field is unsatisfied).

Control plane components can be run on any machine in the cluster, but typically set-up scripts start all *control plane* components on the same machine, and do not run user containers on that machine.

Control Plane – kube-apiserver

- The *API server* exposes the Kubernetes API. The API server is the front end for the Kubernetes *control plane*.

- The main implementation of a Kubernetes API server is *kube-apiserver*.

- *kube-apiserver* is designed to scale horizontally (deploying more instances).

- You can run several instances of *kube-apiserver* and balance traffic between those instances.
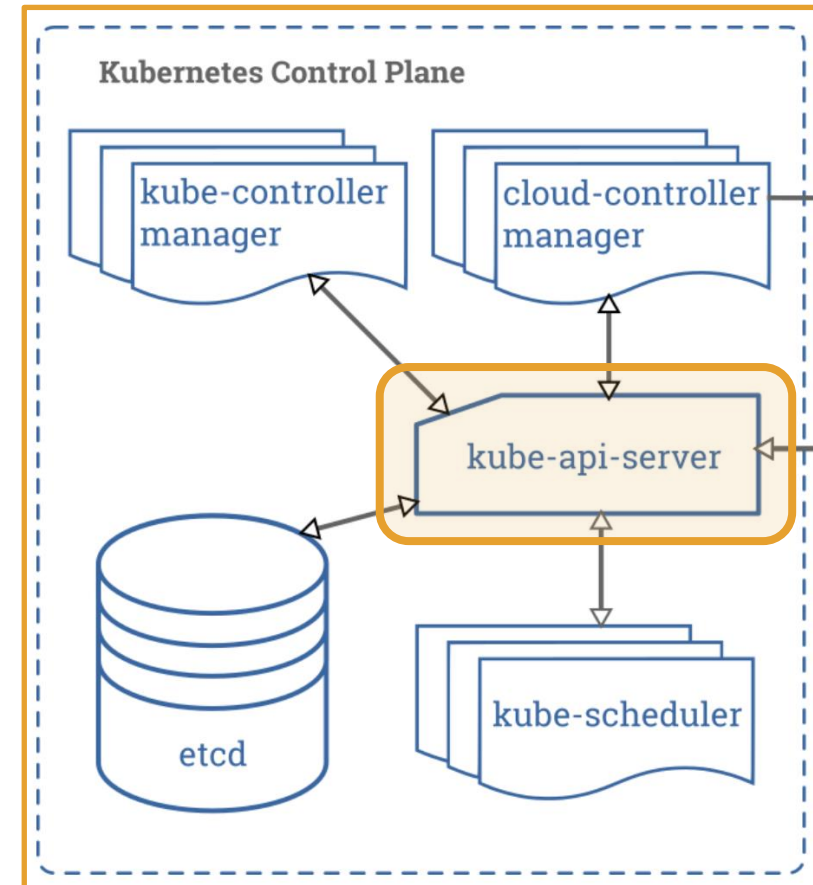
Control Plane – etcd and kube-scheduler

- *Etcd* is a key-value store. It maintains all the *clusters'* data.

- *kube-scheduler* watches for new *Pods* and assigns a *node* to them to run on based on predetermined requirements like:

  - hardware constraints,

  - affinity/anti-affinity specifications,
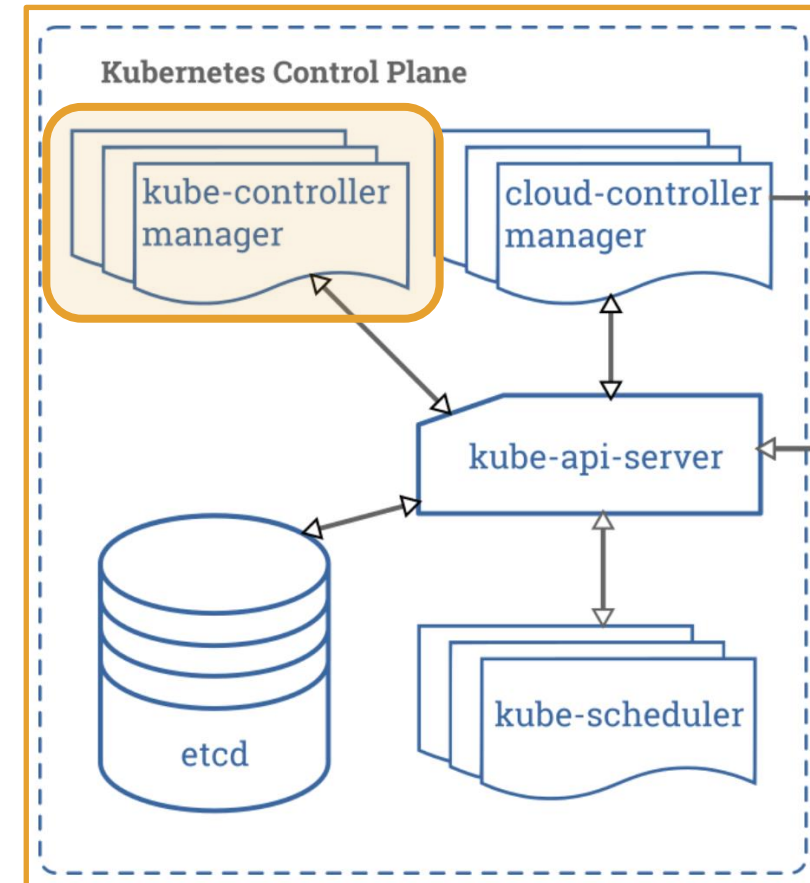
  - deadlines, <mark>and many more</mark>.
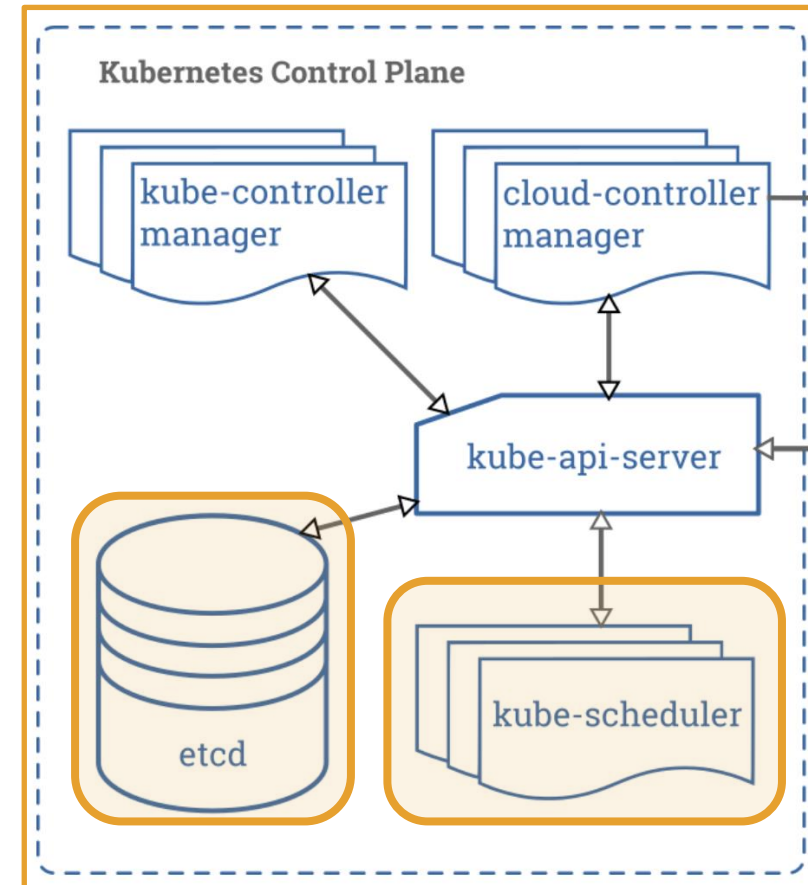
Control Plane – kube-controller-manager

Kube-manager-controller runs the *controller processes*. There are 4 *controllers*:

- Node controller: notices and responds when nodes go down.

- Replication controller: maintains the correct number of pods for every replication controller object in the system.

- Endpoints controller: Populates the Endpoints object (joins Services & Pods).

- Service Account & Token controllers: Create default accounts and API access tokens for new namespaces.
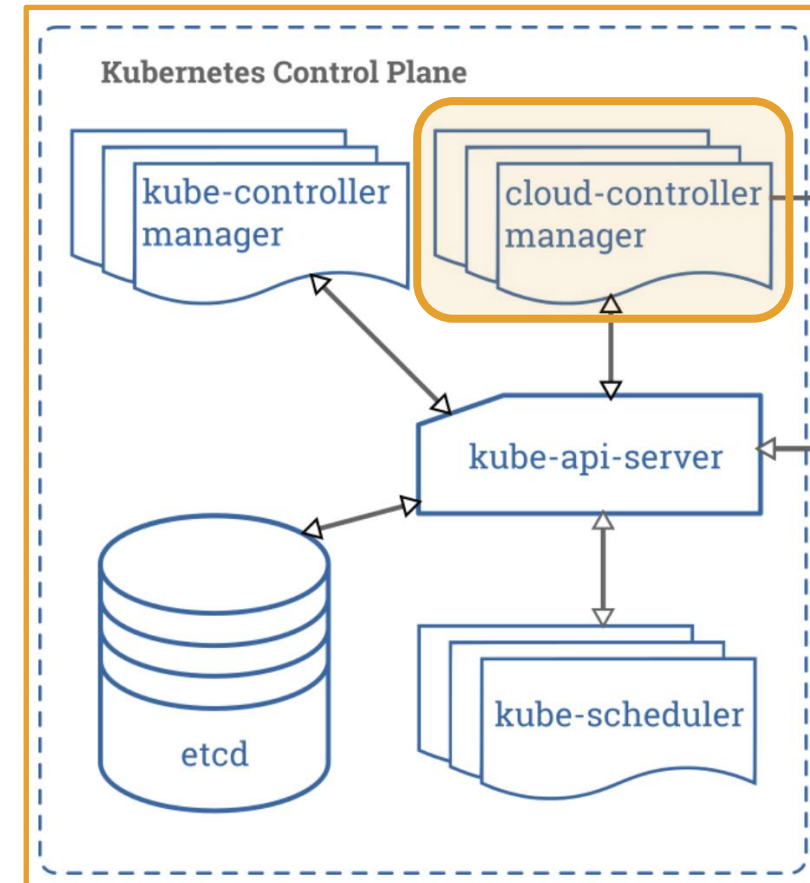
Control Plane – cloud-controller manager (1/2)

https://kubernetes.io/docs/concepts/overview/components/#cloud-controller-manager

The **cloud-controller-manager** allows linking a cluster into the cloud providers API. It will separate the components that interact with the cloud platform from components that only interact with the cluster.

cloud-controller-manager combines several logically independent control loops into a single binary that are run as a single process. Horizontal scaling (running more instances) allows for improved performance or help with failure tolerance.

Control Plane – cloud-controller manager (2/2)

The following controllers can have cloud provider dependencies:

- Node controller: For checking the cloud provider to determine if a node has been deleted in the cloud after it stops responding

- Route controller: For setting up routes in the underlying cloud infrastructure

- Service controller: For creating, updating and deleting cloud provider load balancers.
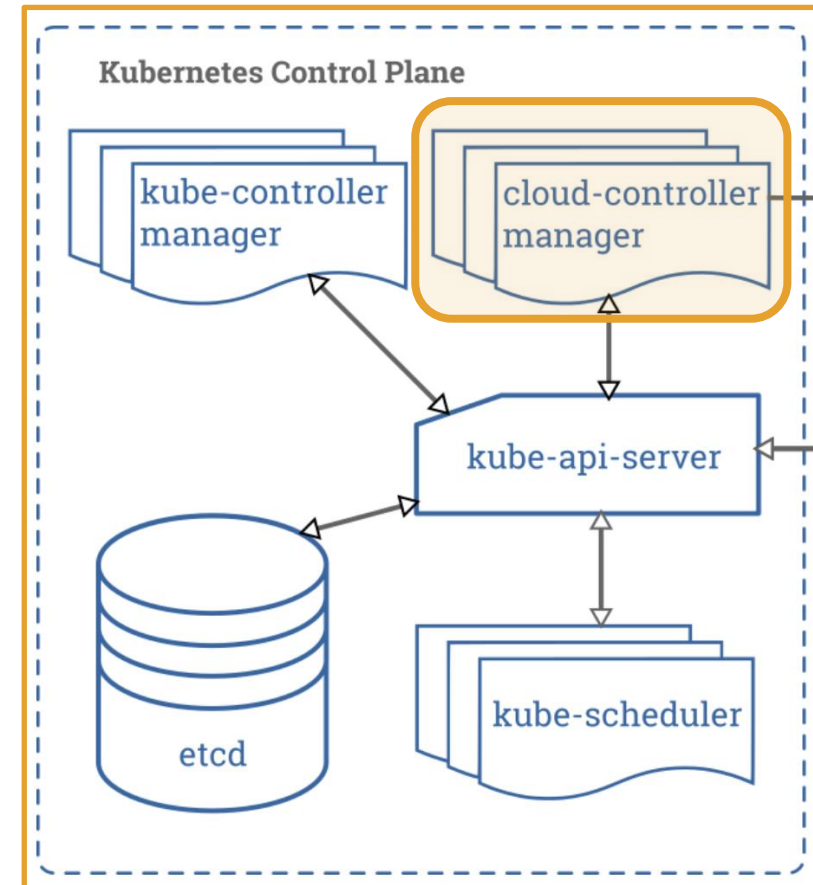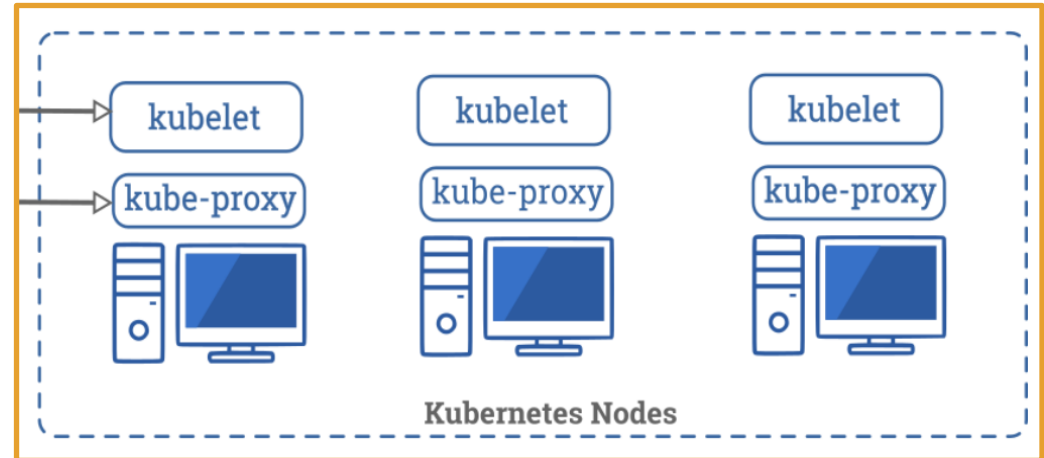
Node Components - Kubelet

A *Kubelet* agent runs on each *node* in the cluster. It is the primary implementer of the *Pod* and Node APIs that drive the container execution layer. The *Kubelet* uses *PodSpecs* to verify that containers described in those *PodSpecs* are running in the *Pods*. The *kubelet* doesn't manage containers which were not created by *Kubernetes*.

Node components run on every *node*, maintain running *pods*, and providing the Kubernetes runtime environment.
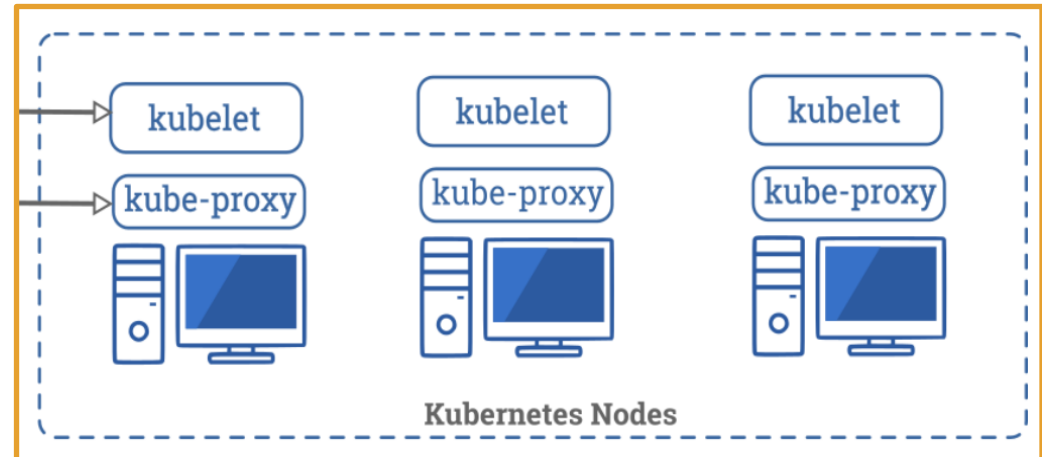
Node Components – kube-proxy

A *kube-proxy* is a network proxy that runs on each *node* in your cluster. *kube-proxy* provides a way to group pods under a common access policy (e.g., *load-balanced*). This creates a virtual IP which clients can access, and which is transparently proxied (forwarded) to the *pods* in a Service. Every *node* runs a *kube-proxy* process. *Kube-proxy* programs IpTables rules to trap access to service IPs and redirect them to the correct backends.

Kubernetes Nodes

Node components run on every *node*, maintain running *pods*, and providing the Kubernetes runtime environment.
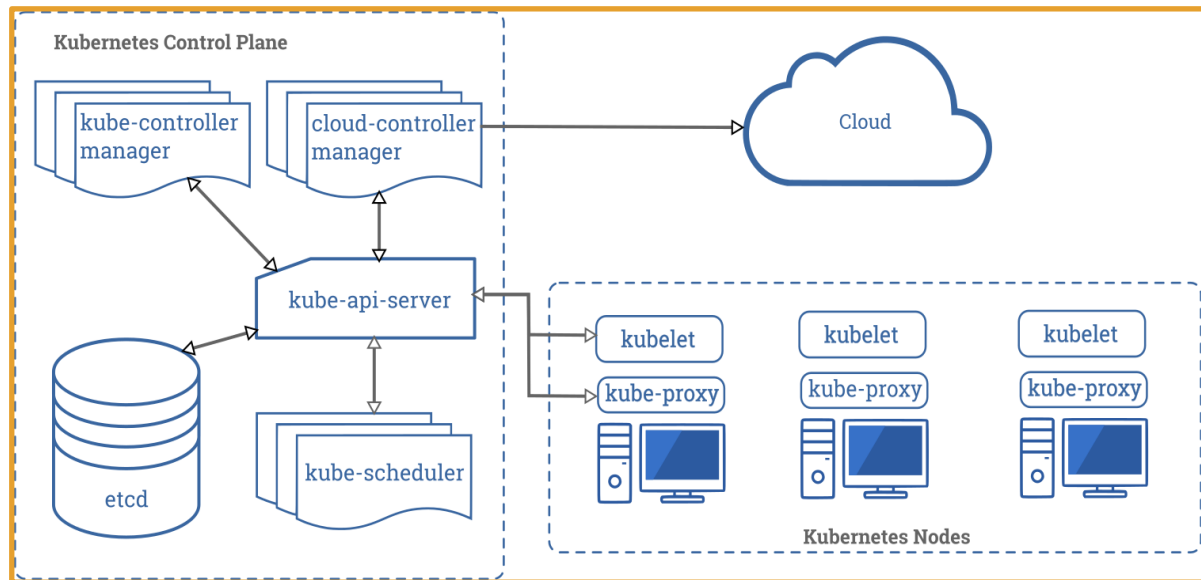
Node – Components

The *container runtime* is the software that is responsible for running containers.

Kubernetes supports several container runtimes.

- Kubernetes Container Runtime Interface (CRI
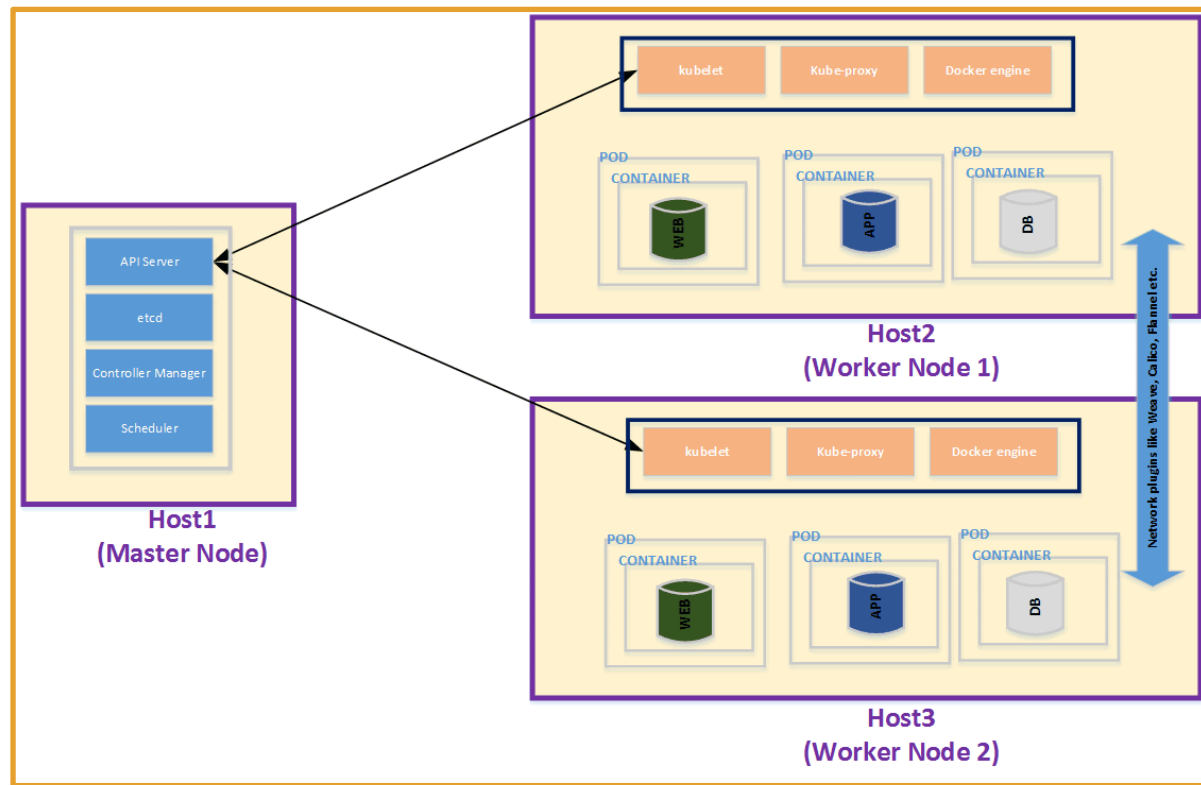- Docker
- containerd
- CRI-O

Node Structure

Each *node* contains the services necessary to run the *Pods* on it, which are managed by the *control plane*.

A *node* may be a virtual or physical machine.

Node Structure

https://kubernetes.io/docs/concepts/architecture/nodes/

A *container* represents a containerized application

The *container* is placed into a *Pod*.

A *Pod* runs on a *Node*.

One or more *Nodes* make up a *Cluster*.

The *Cluster* is the *Kubernetes Workload*.